

Problem:

- Subtask 1: Decode a Base64 BMP, read the blue-channel LSBs, rebuild bytes, extract the number inside ABC{...}.
- Subtask 2: Decode a Base64 PCAP blob, treat it as text, regex-search for ABC{...}, take that big integer, compute (n%10007)+3 (or 0 if not found).
- Subtask 3: Decode Base64 hex bytecode, parse into opcodes, simulate the 8-bit MystVM until HALT, then return the byte at the given memory address.

Assumptions:

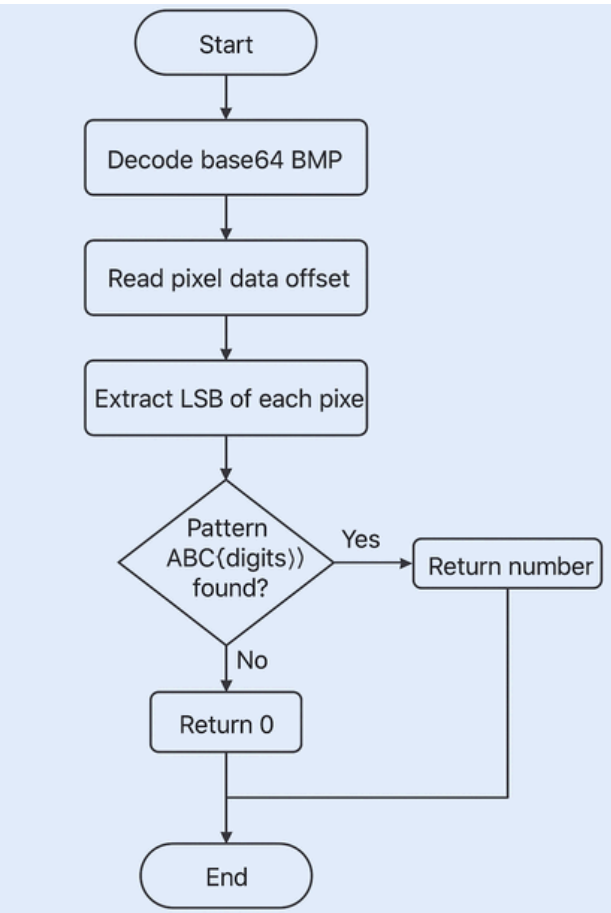
- If any decoding/parsing error occurs, that subtask’s result defaults to 0.
- In case of multiple regex matches, only the first match is used.
- The Base64-decoded PCAP is already reassembled into a single text blob
- Java’s BigInteger can handle the largest possible <number> in the input.
- Registers and memory initialize to zero; %r15 (SP) starts at 0.
- In case of multiple regex matches, only the first match is used.

Subtask 1

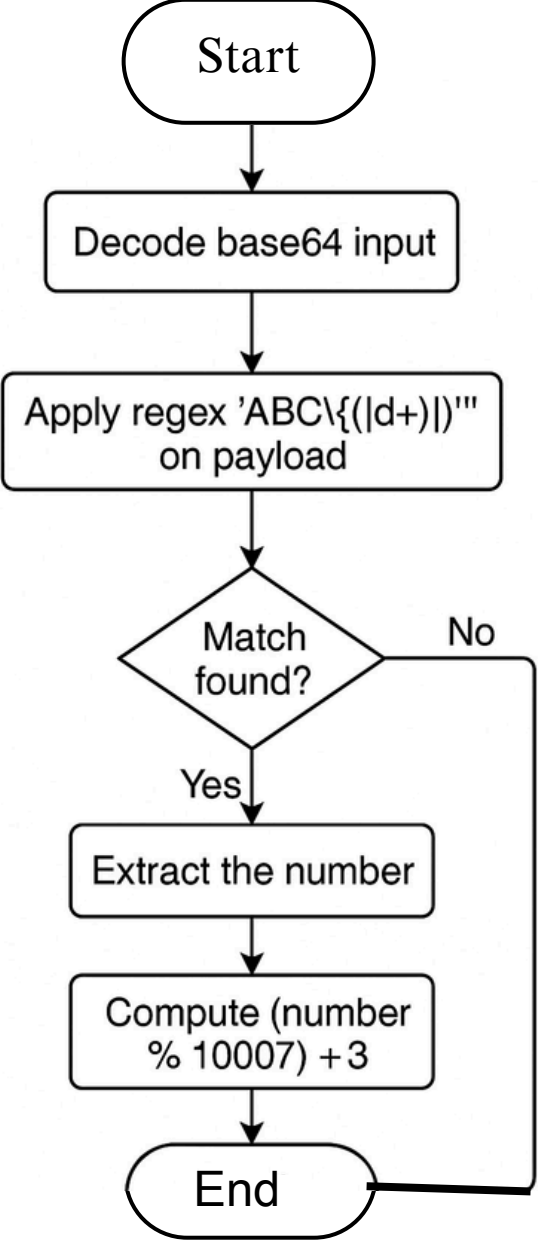
Bitmap File Structure			
Block	Field	Width	Description
BITMAPFILEHEADER Fields: 5 Width: 14 bytes	FileType	2 bytes	A 2 character string value in ASCII. It must be 'BM' or '0x42 0x4D'
	FileSize	4 bytes	An integer (unsigned) representing entire file size in bytes (number of bytes in a BMP image file )
	Reserved	2 bytes	To be utilized by an image processing application. Initialized to '0' integer (unsigned) value.
	Reserved	2 bytes	To be utilized by an image processing application. Initialized to '0' integer (unsigned) value.
	PixelDataOffset	4 bytes	An integer (unsigned) representing the offset of actual pixel data in bytes.

→ Payload

Each pixel: [Blue] [Green] [Red]  
(Little Endian Format)



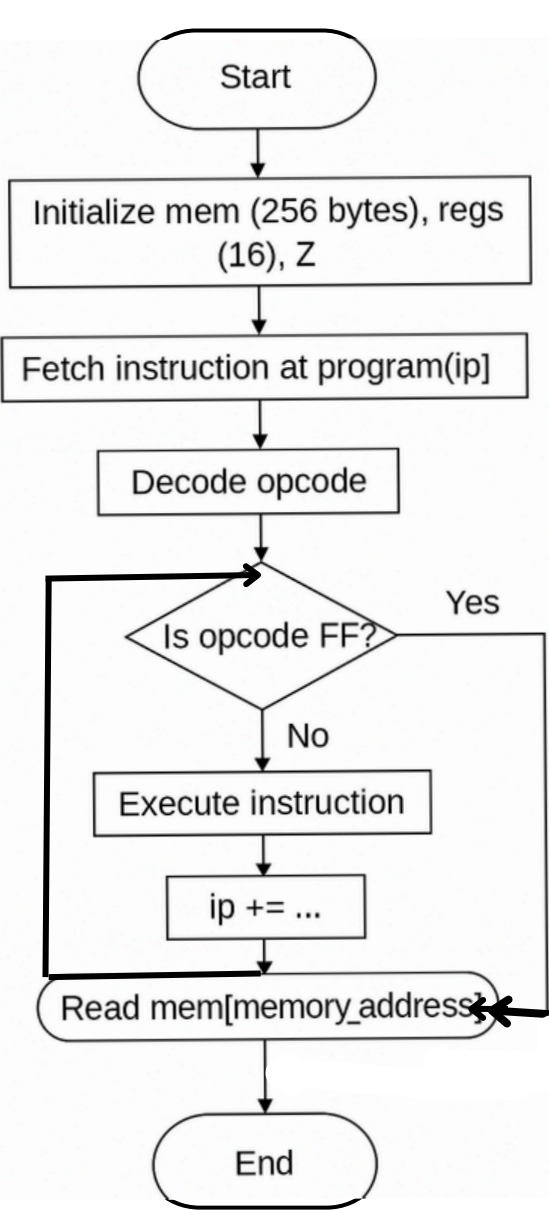
Subtask 2



Why Java?

- Byte-level control & endianness via ByteBuffer and explicit masking (& 0xFF)
- High performance VM loops over millions of ops

Subtask 3



Limitations & improvement

1. No IP/TCP/UDP reassembly

**Curent fix:** Applies regex on raw ISO-8859-1 decoded bytes

**Future fix:** Integrate a pcap library (Scapy/jNetPcap) to reassemble fragments before searching

2. Stack pointer %r15 starts at 0, causing unintended underflow on first call

**Curent fix:** Underflow wraps around silently via & 0xFF

**Future fix:** Initialize %r15 to 0xFF and add bounds checks (error on overflow/underflow).

3. No early exit during BMP LSB scan

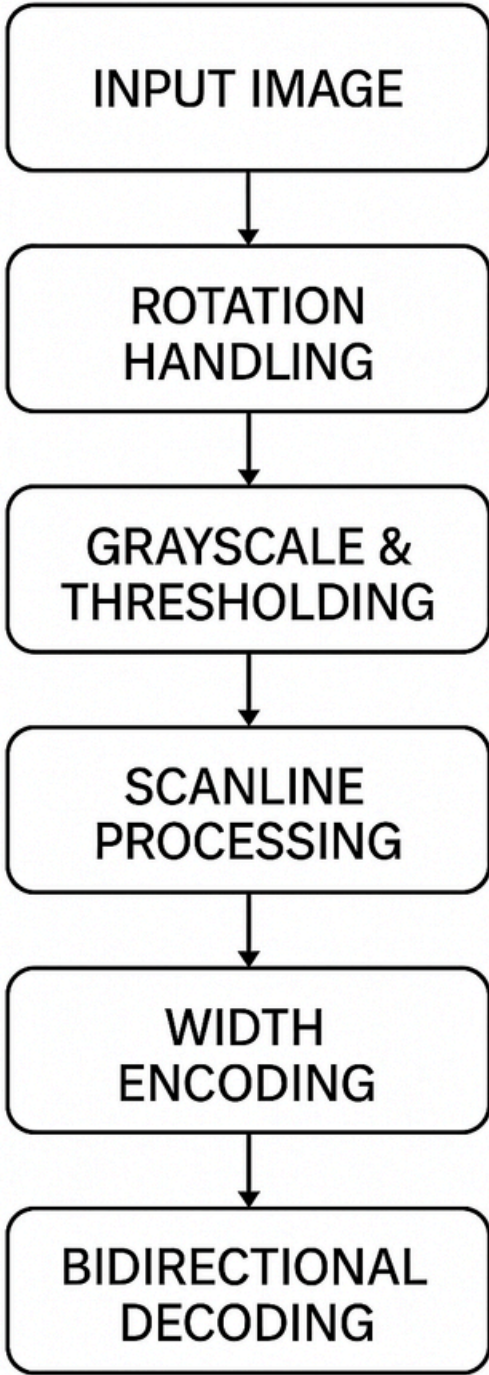
**Curent fix:** Collects full bitstream and reconstructs entire message

**Future fix:** Stop bit extraction as soon as regex match is found to save time


**Problem:** Decode a Code 39 barcode from a noisy, rotated image.

Handle real-world constraints: lighting, skew, inconsistent bar widths.

**Approach:**



**Assumption:** 1 barcode/image



GS2025

Widths (9)	Pattern (N/W)	Char
[3,1,2,1,3,4,5,2,4]	NNNNNWWNW	G
[3,1,4,1,3,2,5,4,2]	NNWNNNWWN	S
[3,1,4,5,3,2,2,1,4]	NNWWNNNNW	2
[3,1,3,5,4,2,4,1,2]	NNNWWNWNN	0
[3,1,4,5,3,2,2,1,4]	NNWWNNNNW	2
[5,1,2,5,4,2,2,1,3]	WNNWWNNNN	5

Why Python?

- Fast prototyping, strong imaging libraries (PIL, NumPy).
- Ideal for iterative logic-heavy pixel processing.

**Limitations & improvement**

**1. Fails under extreme blur/skew**

- Curent fix:** Otsu Thresholding
- Future fix:** Deblurring filters or morphological edge enhancement

**2. Tilt beyond ±5° or vertical barcode**

- Curent fix:** Small-angle rotations handle minor tilts.
- Future fix:** Hough-line based deskew

**3. Corner-placed or tiny barcodes**

- Curent fix:** We check ±3 rows around center
- Future fix:** Sliding-Window + Early Pruning

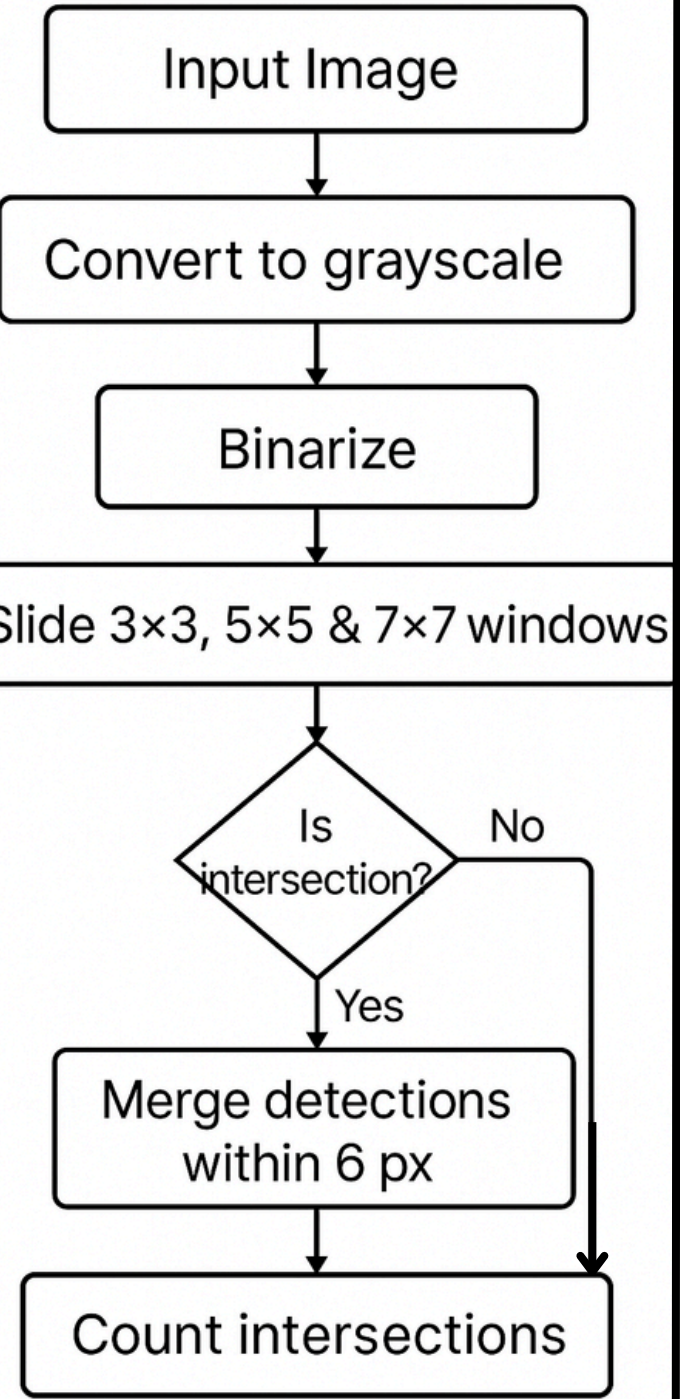
**Otsu’s Thresholding Algorithm**

$$\sigma_B^2 = \frac{[m_G P_1(k) - m(k)]^2}{P_1(k)[1 - P_1(k)]}$$



**Problem:** Count line-line intersections in a 512×512 JPEG graph (axes excluded)

**Approach:**



**3×3 ≥ 3**

0	1	0
0	1	0
1	1	1

Ensures no noise

**5×5 ≥ 6**

0	0	1	0	0
0	0	1	0	0
1	1	1	1	1
0	0	0	0	0
0	0	0	0	0

Confirms the “cross” extends beyond 3×3

**7×7 ≤ 30**

0	0	0	1	1	0	0
0	0	0	1	1	0	0
0	0	0	1	1	0	0
1	1	1	1	1	1	1
0	0	0	1	1	0	0
0	0	0	1	1	0	0

Rejects thick blobs or axes (true 1px or 2px crossings only fill ~13–25 pixels in 7×7)

**Assumptions:**

- Axes are thicker than data lines (≥ 3 px)
- No 2 valid intersection points occur within a 6px radius
- L pattern is not considered as an intersection

**Limitations & improvement**

- 1. Heuristic Threshold Sensitivity**

**Curent fix:** Manually tuned as per 512x512

**Future fix:** Adaptive/local thresholding (Sauvola/Otsu)
- 2. Over-Aggressive Clustering Radius**

**Curent fix:** 6px merge radius to avoid duplicate detections.

**Future fix:** connected-component labeling and centroids or dynamic radius based on local line density.
- 3. Axis & Origin Handling**

**Curent fix:** Implicitly ignored via 3px border skip + 7×7≤30 rule.

**Future fix:** lightweight Hough transform or longest-line removal or padding.

**Alternative Solutions & Generalizations**

- Graph Skeleton & Node Analysis**
- Morphologically thin binary image to a 1-px skeleton
  - Treat skeleton as a graph: pixels = nodes, adjacencies = edges
  - Identify graph nodes of degree ≥ 3 as intersections
  - Scales to complex diagrams or network maps